

Members: Jonathan Ng, Melody Song, Sunny Kim, Van Vai Vivien Lee

Database: PostgreSQL

Q1. Stored Procedure

This function applies a discount to bookings that exceed a certain price . There are two float input arguments-- 1) discount percentage; 2) minimum price to qualify for the booking.

Stored Procedure

```
CREATE OR REPLACE FUNCTION discount(IN percent float, IN min_price float)
RETURNS void AS $$
BEGIN
    DECLARE old_book VARCHAR(20);
    DECLARE old_price float;
    DECLARE at_end INT DEFAULT 0;

    DECLARE curs1 CURSOR FOR SELECT bookingid, price FROM Booking;

    BEGIN
        OPEN curs1;
        FETCH curs1 INTO old_book, old_price;
        WHILE at_end = 0 LOOP
            IF (old_price > min_price) THEN
                UPDATE booking SET price = old_price * percent WHERE booking.bookingid = old_book;
            END IF;
            FETCH curs1 INTO old_book, old_price;
            EXIT WHEN NOT FOUND;
        END LOOP;
        CLOSE curs1;
    END;
END;
$$ LANGUAGE plpgsql;
```

Before the function discount is called, the records of booking are as below:

```
cs421=> SELECT * FROM booking;
```

bookingid	traveldate	deptlocation	arrvlocation	price	status	bookingtype	disclaimer	noofcustomers
b400	2020-03-29	Montreal, Canada	Amsterdam, Netherlands	2700	Pending	true	none	2
b201	2020-03-01	Toronto, Canada	Beijing, China	3700	Incomplete	false	none	2
b066	2020-04-08	London, England	Reykjavik, Iceland	2280	Incomplete	true	none	1
b449	2020-04-25	Orlando, United States	Paris, France	2210	Pending	false	none	2
b370	2020-04-05	Ottawa, Canada	Calgary, Canada	2500	Successful	true	none	3
b173	2020-04-29	New York, United States	Rome, Italy	3300	Successful	true	none	3
b437	2020-03-15	Seoul, South Korea	Toronto, Canada	3550	Successful	false	none	1
b279	2020-04-29	Mexico City, Mexico	Austin, United States	1825	Successful	true	none	2
b170	2020-04-17	Berlin, Germany	Chicago, United States	2580	Successful	false	none	3
b442	2020-03-29	Kyoto, Japan	Las Vegas, United States	2830	Successful	true	none	1
b222	2020-03-29	Kyoto, Japan	Beijing, China	800	Successful	true	none	1

(11 rows)

Now we test this function passing 0.8 (80% of original price) and 1000 as arguments:

```
cs421=> SELECT discount(0.8,1000);
discount
-----
(1 row)
```

The output records show that bookings with original price > 1000 got a 20% discount

```
cs421=> SELECT * FROM Booking;
```

bookingid	traveldate	deptlocation	arrvlocation	price	status	bookingtype	disclaimer	noofcustom
b400	2020-03-29	Montreal, Canada	Amsterdam, Netherlands	2160	Pending	true	none	
b201	2020-03-01	Toronto, Canada	Beijing, China	2960	Incomplete	false	none	
b066	2020-04-08	London, England	Reykjavik, Iceland	1824	Incomplete	true	none	
b449	2020-04-25	Orlando, United States	Paris, France	1768	Pending	false	none	
b370	2020-04-05	Ottawa, Canada	Calgary, Canada	2000	Successful	true	none	
b173	2020-04-29	New York, United States	Rome, Italy	2640	Successful	true	none	
b437	2020-03-15	Seoul, South Korea	Toronto, Canada	2840	Successful	false	none	
b279	2020-04-29	Mexico City, Mexico	Austin, United States	1460	Successful	true	none	
b170	2020-04-17	Berlin, Germany	Chicago, United States	2064	Successful	false	none	
b442	2020-03-29	Kyoto, Japan	Las Vegas, United States	2264	Successful	true	none	
b222	2020-03-29	Kyoto, Japan	Beijing, China	800	Successful	true	none	

2. Java Application

1) Select a single attribute from a table

Choice 1: Select an attribute from a table

Number of ...

Choice 3: Insert new data or delete from table.
Syntax: INSERT INTO table (cols-opt) VALUES params or DELETE FROM table WHERE cols IN (param)

associates

consultatio...

bookingid

b201

b449

b279

b170

b222

b442

b370

b173

b437

b066

b400

By inputting the table name and parameters to the fields, we can acquire a single column from any table

2) We find customers without insurance by doing two parameters except the other

```
SELECT customer.custid FROM customer
EXCEPT
SELECT associates.custid FROM associates;
```

3) Inserting & Deleting a new record to & from a table

Insert: Using employee as an example, inserting a new employee record

- Before insertion (only displayed full records, i.e. records don't have null values)

Number of ...

Choice 3: Insert new data or delete from table.
Syntax: INSERT INTO table (cols-opt) VALUES params or DELETE FROM table WHERE cols IN (param)

employee

eid

'e9897'

Insert

Delete

associates

consultatio...

e159	McLennan Roland
e132	Pascal Siakam
e202	Will Wellington
e095	Prakash Panangaden
e050	Juliet Romeo
e173	Carl Gauss
e060	Alan Turing
e105	Albert Eggstein
e033	Mitzi Fitzgerald
e066	Love Blind
e111	tester
c111	tets
c123	parametersasd
e99000	david davidson
e1123456	Harley Quinn
e556776	Test mcTest
e1234567	Test mcTester
e6666666	test test
e890890	peter johnson
e214134134	a
e9019304	lucy
e1356	etest
e896868	tom tomson
e12345678	dad dad
e123456789	null
e3232	null
e12345624	null
e9887	null

- After insertion: we inserted eid = 'e9897'

GUI Framework: Fill out the necessary information and select the choice button
Press Quit to close the connection and exit the GUI framework

Quit

Choice 1: Select an attribute from a table

eid

Choice 1

Number of c...

Choice 3: Insert new data into a table.

employee

eid

'e9897'

Insert

Delete

Choice 4

Choice 5

e159 McLennan Roland
e132 Pascal Siakam
e202 Will Wellington
e095 Prakash Panangaden
e050 Juliet Romeo
e173 Carl Gauss
e060 Alan Turing
e105 Albert Eggstein
e033 Mitzi Fitzgerald
e066 Love Blind
e111 tester
c111 tets
c123 parametersasd
e99000 david davidson
e1123456 Harley Quinn
e556776 Test mcTest
e1234567 Test mcTester
e666666 test test
e890890 peter johnson
e214134134 a
e9019304 lucy
e1356 etest
e896868 tom tomson
e12345678 dad dad
e123456789 null
e3232 null
e12345624 null
e9887 null
e9897 null

5) Counting the # of consultations each employee has

```
SELECT eid, COUNT(sessionid) FROM  
(SELECT employee.eid, consults.sessionid FROM employee  
LEFT JOIN consults on  
employee.eid = consults.eid) as A  
GROUP BY A.eid;
```

GUI Framework: Fill out the necessary information and select the choice button
Press Quit to close the connection and exit the GUI framework

Choice 1: Select an attribute from a table

Attribute Name Table Name

Choice 3: Insert new data or delete from table.
Syntax: INSERT INTO table (cols-opt) VALUES params or DELETE FROM table WHERE cols IN (param)

Table name imns (optional) Parameters

eid	count(sessionid)
e214134134	0
e3232	0
e202	1
e033	1
e12345624	0
e890890	0
e159	1
c111	0
e132	1
e1234567	0
e9019304	0
e123456789	0
c123	0
e060	0
e1356	0
e99000	0
e6666666	0
e105	1
e173	1
e066	1
e095	1
e111	0
e12345678	0
e896868	0
e1123456	0
e556776	0
e050	1

3. Creating indices

**Note: Unlike DB2, postgresSQL does not have the option to create a clustered index in the index definition.*

a. Index on attribute status of relation booking (clustered)

- The primary key of booking is bookingid

This index is useful because a travel agency will oftentimes have to track the status of a booking and see if further actions need to be executed. To recall, there are three possible values for booking status-- "pending", "incomplete" and "successful". By using an index on this attribute we can quickly select pages with matching records.

For example, if we do `SELECT * FROM Booking WHERE status = "successful"`. Instead of having to scan all the data pages, we can read one leaf page (possibly more if it spreads over more than one leaf node) and corresponding data pages. Since we have a clustered index, assuming we have 30% matching records, we will only have to read 30% of the data pages to retrieve all the records. The total IO cost from reading will go down from 100% data pages to 30% data pages plus one leaf page.

```
CREATE INDEX booking_status ON Booking(status);  
CLUSTER booking USING booking_status;
```

b. Index on attribute eid of relation consults (unclustered)

- The primary key of consults is (customer id, eid)

This index will be useful if the travel agency wants to see how many consultation sessions a particular employee has conducted have conducted.

For example, if we do `SELECT COUNT(sessionid) FROM consults WHERE eid = @xyz`. If we have the index on eid (employee id), Instead of having to scan all the data pages in the relation, we can read one leaf page and corresponding data pages only. We will use a clustered index so all records belonging to the same employee are next to each other. This can bring our IO cost down from reading all data pages to possibly 1 leaf page and say only 1 data page (assuming all the matching records of this employee can fit into one data page).

```
CREATE INDEX consults_eid ON consults(eid);  
CLUSTER consults USING consults eid;
```


b. Average price of hotel per night by location

1) SQL Query to generate data (relation Hotel)

cs421=> \COPY (SELECT* FROM hotel) TO hotel.csv WITH CSV

2) Excel pivot chart image



3) Excel file

	A	B	C	D	E	F	G	H	I
1	hotelid	name	location	checkindate	checkoutdate	price	roomtype	num_days	price_per_day
2	h647	Minuit Inns	Paris, France	2020-04-25	2020-05-12	1200	Comfort Plus	17	70.58823529
3	h123	West Bestern	Calgary, Canada	2020-04-05	2020-04-19	2000	Superior	14	142.8571429
4	h234	Holland Inns	Amsterdam, Netherlands	2020-03-29	2020-04-11	1800	Deluxe Double	13	138.4615385
5	h788	Stay Here Please	Beijing, China	2020-03-01	2020-03-09	2500	Executive Suite	8	312.5
6	h532	La Suite Bella	Rome, Italy	2020-04-29	2020-05-10	2500	Deluxe Single	11	227.2727273
7	h677	West Bestern	Toronto, Canada	2020-03-15	2020-03-31	2250	Superior	16	140.625
8	h111	Hotel Motel	Austin, United States	2020-04-29	2020-05-13	1425	Regular	14	101.7857143
9	h093	Fishy Suites	Reykjavik, Iceland	2020-04-08	2020-04-19	1600	Comfort	11	145.4545455
10	h543	Hotel Motel	Chicago, United States	2020-04-17	2020-04-27	1700	Garden view	10	170
11	h948	Hotel Motel	Las Vegas, United States	2020-03-29	2020-04-12	1840	Ocean view	14	131.4285714
12									
13									

*The last two columns- num_days and price_per_day are additional manipulations done in excel

5. Creativity points

a. Additional stored procedure

This function takes in nothing as an input but computes the sum (aggregate) of prices of each booking in the Booking relation. This is important to help the business learn how much sales they are generating in total.

```
CREATE OR REPLACE FUNCTION sum_booking_price() RETURNS float AS $total$  
DECLARE total float;  
BEGIN  
    SELECT SUM(Booking.price) into total FROM Booking;  
    RETURN total;  
END;  
$total$ LANGUAGE plpgsql;
```